

Important Components of the Python Scientific Stack

Continuum Analytics Anaconda

Anaconda, a free product of Continuum Analytics (www.continuum.io), is a virtually complete scientific stack for Python. It includes both the core Python interpreter and standard libraries as well as most modules required for data analysis. Anaconda is free to use and modules for accelerating the performance of linear algebra on Intel processors using the Math Kernel Library (MKL) are available (free to academic users and for a small cost to non-academic users).

Continuum Analytics also provides other high-performance modules for reading large data files or using the GPU to further accelerate performance for an additional, modest charge.

Installing Anaconda

Most importantly, installation is extraordinarily easy on Windows, Linux and OS X. Anaconda is also simple to update to the latest version using `conda update conda` `conda update anaconda`

NumPy

NumPy provides a set of array and matrix data types which are essential for statistics, econometrics and data analysis.

SciPy contains a large number of routines needed for analysis of data.

The most important include a wide range of random number generators, linear algebra routines and optimizers.

SciPy depends on NumPy.

IPython



IPython provides an interactive Python environment which enhances productivity when developing code or performing interactive data analysis.

matplotlib and seaborn

- ▶ matplotlib provides a plotting environment for 2D plots, with limited support for 3D plotting.
- ▶ seaborn is a Python package that improves the default appearance of matplotlib plots without any additional code.

Conference Mission

PyData is a gathering of users and developers of data analysis tools in Python. The goals are to provide Python enthusiasts a place to share ideas and learn from each other about how best to apply our language and tools to ever-evolving challenges in the vast realm of data management, processing, analytics, and visualization.

We aim to be an accessible, community-driven conference, with tutorials for novices, advanced topical workshops for practitioners, and opportunities for package developers and users to meet in person.

A major goal of the conference is to provide a venue for users across all the various domains of data analysis to share their experiences and their techniques, as well as highlight the triumphs and potential pitfalls of using Python for certain kinds of problems.

pandas

- ▶ pandas provides high-performance data structures.

pandas



pandas is a high-performance module that provides a comprehensive set of structures for working with data. *pandas* excels at handling structured data, such as data sets containing many variables, working with missing values and merging across multiple data sets.

pandas

While extremely useful, *pandas* is not an essential component of the Python scientific stack unlike NumPy, SciPy or matplotlib, and so while *pandas* doesn't make data analysis possible in Python, it makes it much easier. *pandas* also provides high-performance, robust methods for importing from and exporting to a wide range of formats.

Performance Modules : Cython and Numba

A number of modules are available to help with performance. These include Cython and Numba. Cython is a Python module which facilitates using a simple Python-derived creole to write functions that can be compiled to native (C code) Python extensions.

Numba uses a method of just-in-time compilation to translate a subset of Python to native code using Low-Level VirtualMachine (LLVM).