

Important Components of the Python Scientific Stack

Continuum Analytics Anaconda

Anaconda, a free product of Continuum Analytics (www.continuum.io), is a virtually complete scientific stack for Python. It includes both the core Python interpreter and standard libraries as well as most modules required for data analysis. Anaconda is free to use and modules for accelerating the performance of linear algebra on Intel processors using the Math Kernel Library (MKL) are available (free to academic users and for a small cost to non-academic users).

Continuum Analytics also provides other high-performance modules for reading large data files or using the GPU to further accelerate performance for an additional, modest charge.

Installing Anaconda

Most importantly, installation is extraordinarily easy on Windows, Linux and OS X. Anaconda is also simple to update to the latest version using `conda update conda conda update anaconda`

NumPy

NumPy provides a set of array and matrix data types which are essential for statistics, econometrics and data analysis.

SciPy contains a large number of routines needed for analysis of data.

The most important include a wide range of random number generators, linear algebra routines and optimizers.

SciPy depends on NumPy.

IPython

IPython provides an interactive Python environment which enhances productivity when developing code or performing interactive data analysis.

matplotlib and seaborn

- ▶ matplotlib provides a plotting environment for 2D plots, with limited support for 3D plotting.
- ▶ seaborn is a Python package that improves the default appearance of matplotlib plots without any additional code.

pandas

- ▶ pandas provides high-performance data structures.

pandas

pandas is a high-performance module that provides a comprehensive set of structures for working with data. *pandas* excels at handling structured data, such as data sets containing many variables, working with missing values and merging across multiple data sets.

pandas

While extremely useful, *pandas* is not an essential component of the Python scientific stack unlike NumPy, SciPy or matplotlib, and so while *pandas* doesn't make data analysis possible in Python, it makes it much easier. *pandas* also provides high-performance, robust methods for importing from and exporting to a wide range of formats.

Performance Modules : Cython and Numba

A number of modules are available to help with performance. These include Cython and Numba. Cython is a Python module which facilitates using a simple Python-derived creole to write functions that can be compiled to native (C code) Python extensions.

Numba uses a method of just-in-time compilation to translate a subset of Python to native code using Low-Level VirtualMachine (LLVM).

Versions of Python

- ▶ Version 2.7
- ▶ Version 3

Python Coding Conventions

There are a number of common practices which can be adopted to produce Python code which looks more like code found in other modules:

- ▶ Use 4 spaces to indent blocks avoid using tab, except when an editor automatically converts tabs to 4 spaces
- ▶ Avoid more than 4 levels of nesting, if possible
- ▶ Limit lines to 79 characters. The `\` symbol can be used to break long lines
- ▶ Use two blank lines to separate functions, and one to separate logical sections in a function.

Python Coding Conventions

- ▶ Use ASCII mode in text editors, not UTF-8
- ▶ One module per import line
- ▶ Avoid from module import * (for any module). Use either from module import func1, func2 or import module as shortname.
- ▶ Follow the NumPy guidelines for documenting functions

More suggestions can be found in PEP8.

Part 2 : Other Interesting Python Packages

statsmodels

- ▶ statsmodels provides a large range of cross-sectional models aswell as time-series models.
- ▶ statsmodels uses a model descriptive language (provided via the Python package patsy) to formulate the model when working with pandas DataFrames.
- ▶ Models supported include linear regression, generalized linear models, limited dependent variable models, ARMA and VAR models.

pytz and babel

pytz and babel provide extended support for time zones and formatting information.

rpy2 provides an interface for calling R 3.0.x in Python, as well as facilities for easily moving data between the two platforms.

PyTables and h5py

PyTables and h5py both provide access to HDF5 files, a flexible data storage format optimized for numeric data.

subfiles framed amsmath amssymb

Data Structures

pandas introduces two new data structures to Python - **Series** and **DataFrame**, both of which are built on top of NumPy.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
pd.set_option('max_columns', 50)
```

Series

Series is a one-dimensional labeled array capable of holding any data type (integers, strings, floating point numbers, Python objects, etc.). The axis labels are collectively referred to as the index. The basic method to create a Series is to call:

```
s = Series(data, index=index)
```

Here, data can be many different things:

- ▶ a Python dict
- ▶ an ndarray
- ▶ a scalar value (like 5)

- ▶ A Series is a one-dimensional object similar to an array, list, or column in a table.
- ▶ It will assign a labeled index to each item in the Series.
- ▶ By default, each item will receive an index label from 0 to N, where N is the length of the Series minus one.

```
# create a Series with an arbitrary list
s = pd.Series([7, 'Heisenberg', 3.14, -1789710578,
               'Happy Eating!'])
s
```

Series

Output from Previous Slide

```
0          7
1    Heisenberg
2        3.14
3   -1789710578
4   Happy Eating!
dtype: object
```

Alternatively, you can specify an index to use when creating the Series.

```
s = pd.Series([7, 'Heisenberg', 3.14, -1789710578,
               'Happy Eating!'],
               index=['A', 'Z', 'C', 'Y', 'E'])
s
```

```
A          7
Z    Heisenberg
C          3.14
Y    -1789710578
E    Happy Eating!
dtype: object
```

Series

The Series constructor can convert a dictionary as well, using the keys of the dictionary as its index.

```
d = {'Chicago': 1000, 'New York': 1300, 'Portland': 900,  
     'Austin': 450, 'Boston': None}  
cities = pd.Series(d)  
cities  
Out[4]:  
Austin          450  
Boston         NaN  
Chicago        1000  
New York       1300  
Portland        900  
San Francisco  1100  
dtype: float64
```

Series

You can use the index to select specific items from the Series ...

```
cities['Chicago']  
Out[5]:  
1000.0
```

Series

```
cities[['Chicago', 'Portland', 'San Francisco']]
Out[6]:
Chicago          1000
Portland          900
San Francisco    1100
dtype: float64
```

Series

You can use **boolean indexing** for selection.

```
cities[cities < 1000]
```

```
Out[7]:
```

```
Austin      450
```

```
Portland    900
```

```
dtype: float64
```

That last one might be a little strange, so let's make it more clear
- `cities < 1000` returns a Series of True/False values, which
we then pass to our Series `cities`, returning the corresponding True
items.


```
less_than_1000 = cities < 1000
print less_than_1000
print '\n'
print cities[less_than_1000]
Austin           True
Boston           False
Chicago          False
New York         False
Portland         True
San Francisco    False
dtype: bool
```

```
Austin      450
Portland    900
dtype: float64
```

You can also change the values in a Series on the fly.

```
# changing based on the index

print 'Old value:', cities['Chicago']

cities['Chicago'] = 1400
print 'New value:', cities['Chicago']

Old value: 1000.0
New value: 1400.0
```

Changing values using boolean logic

```
print cities[cities < 1000]
print '\n'
cities[cities < 1000] = 750
```

```
print cities[cities < 1000]
Austin      450
Portland    900
dtype: float64
```

```
Austin      750
Portland    750
dtype: float64
```

Working with Series

What if you aren't sure whether an item is in the Series? You can check using idiomatic Python.

```
print 'Seattle' in cities
print 'San Francisco' in cities
False
True
```

Mathematical operations can be done using scalars and functions.

```
# divide city values by 3
cities / 3
Out[12]:
Austin          250.000000
Boston          NaN
Chicago         466.666667
New York        433.333333
Portland        250.000000
San Francisco   366.666667
dtype: float64
```

```
# square city values
np.square(cities)
Out[13]:
Austin          562500
Boston          NaN
Chicago         1960000
New York        1690000
Portland        562500
San Francisco   1210000
dtype: float64
```

You can add two Series together, which returns a union of the two Series with the addition occurring on the shared index values. Values on either Series that did not have a shared index will produce a NULL/NaN (not a number).

```
print cities[['Chicago', 'New York', 'Portland']]
print'\n'
print cities[['Austin', 'New York']]
print'\n'
print cities[['Chicago', 'New York', 'Portland']] + cities[['Austin', 'New York']]
```

Chicago	1400
New York	1300
Portland	750

dtype: float64

Austin	750
New York	1300

dtype: float64

Austin	NaN
Chicago	NaN
New York	2600
Portland	NaN

dtype: float64

Working with Series

NULL Checking

- ▶ Notice that because Austin, Chicago, and Portland were not found in both Series, they were returned with NULL/NaN values.
- ▶ NULL checking can be performed with `isnull()` and `notnull()`.

Return a boolean series indicating which values aren't NULL

```
cities.notnull()
```

Austin	True
Boston	False
Chicago	True
New York	True
Portland	True
San Francisco	True
dtype:	bool

Using boolean logic to grab the NULL cities

```
print cities.isnull()
print '\n'
print cities[cities.isnull()]
Austin           False
Boston           True
Chicago          False
New York         False
Portland         False
San Francisco    False
dtype: bool

Boston    NaN
dtype: float64
```

Special Arrays

Functions are available to construct a number of useful, frequently encountered arrays.

ones

`ones` generates an array of 1s and is generally called with one argument, a tuple, containing the size of each dimension. `ones` takes an optional second argument (`dtype`) to specify the data type. If omitted, the data type is `float`.

```
>>> M, N = 5, 5
>>> x = ones((M,N)) # M by N array of 1s
>>> x = ones((M,M,N)) # 3D array
>>> x = ones((M,N), dtype=int32) # 32bit integers
```

zeros

`zeros` produces an array of 0s in the same way `ones` produces an array of 1s, and commonly used to initialize an array to hold values generated by another procedure. `zeros` takes an optional second argument (`dtype`) to specify the data type. If omitted, the data type is float.

```
>>> x = zeros((M,N)) # M by N array of 0s
>>> x = zeros((M,M,N)) # 3D array of 0s
>>> x = zeros((M,N),dtype=int64) # 64 bit integers
```

ones

`ones_like` creates an array with the same shape and data type as the input. Calling `ones_like(x)` is equivalent to calling `ones(x.shape,x.dtype)`. `zeros_like` creates an array with the same size and shape as the input. Calling `zeros_like(x)` is equivalent to calling `zeros(x.shape,x.dtype)`.

empty

`empty` produces an empty (uninitialized) array to hold values generated by another procedure. `empty` takes an optional second argument (`dtype`) which specifies the data type. If omitted, the data type is float.

```
>>> x = empty((M,N)) # M by N empty array  
>>> x = empty((N,N,N,N)) # 4D empty array  
>>> x = empty((M,N),dtype=float32) # 32bit
```


floats (single precision)

- ▶ Using `empty` is slightly faster than calling `zeros` since it does not assign 0 to all elements of the array the empty array created will be populated with (essential random) non-zero values.
- ▶ `empty_like` creates an array with the same size and shape as the input.
- ▶ Calling `empty_like(x)` is equivalent to calling `empty(x.shape,x.dtype)`.

eye, identity

`eye` generates an identity array an array with ones on the diagonal, zeros everywhere else. Normally, an identity array is square and so usually only 1 input is required. More complex zero-padded arrays containing an identity matrix can be produced using optional inputs.

```
>>> In = eye(N)
```

`identity` is a virtually identical function with similar use, `In = identity(N)`.

The Normal Distribution - normal

The main commands

- ▶ `normal()` generates a set of random numbers from a standard Normal distribution.
- ▶ `normal(mu, sigma)` generates draws from a Normal distribution with mean μ and standard deviation σ .
- ▶ `normal(mu, sigma, (10,10))` generates a 10 by 10 array of draws from a Normal with mean μ and standard deviation σ .
- ▶ `normal(mu, sigma)` is equivalent to `mu + sigma * standard_normal()`.

The Poisson Distribution - poisson

- ▶ `poisson()` generates a set of random numbers from a Poisson distribution with $\lambda = 1$.
- ▶ `poisson(lambda)` generates a draw from a Poisson distribution with expectation λ .
- ▶ `poisson(lambda, (10,10))` generates a 10 by 10 array of draws from a Poisson distribution with expectation λ .

standard_t

`standard_t(nu)` generates a set of random numbers from a Students t with shape parameter ν .

`standard_t(nu, (10,10))` generates a 10 by 10 array of draws from a Students t with shape parameter ν .

uniform

`uniform()` generates a uniform random variable on $(0, 1)$.

`uniform(low, high)` generates a uniform on (l, h) .

`uniform(low, high, (10,10))` generates a 10 by 10 array of uniforms on (l, h) .

Continuous Random Variables

SciPy contains a large number of functions for working with continuous random variables. Each function resides in its own class (e.g. `norm` for Normal or `gamma` for Gamma), and classes expose methods for random number generation, computing the PDF, CDF and inverse CDF, fitting parameters using MLE, and computing various moments. The methods are listed below, where `dist` is a generic placeholder for the distribution name in SciPy.

► `dist.rvs`

Pseudo-random number generation. Generically, `rvs` is called using `dist.rvs(*args, loc=0, scale=1, size=size)` where `size` is an n-element tuple containing the size of the array to be generated.

► `dist.pdf`

Probability density function evaluation for an array of data (element-by-element). Generically, `pdf` is called using `dist.pdf(x, *args, loc=0, scale=1)` where `x` is an array that contains the values to use when evaluating PDF.

► `dist.cdf`

Cumulative distribution function evaluation for an array of data (element-by-element). Generically, `cdf` is called using `dist.cdf(x, *args, loc=0, scale=1)` where `x` is an array that contains the values to use when evaluating CDF.

► `dist.ppf`

Inverse CDF evaluation (also known as percent point function) for an array of values between 0 and 1. Generically, `ppf` is called using `dist.ppf(p, *args, loc=0, scale=1)` where `p` is an array with all elements between 0 and 1 that contains the values to use when evaluating inverse CDF.

► `dist.fit`

Estimate shape, location, and scale parameters from data by maximum likelihood using an array of data.

Generically, fit is called using `dist.fit(data, *args, floc=0, fscale=1)` where data is a data array used to estimate the parameters.

`floc` forces the location to a particular value (e.g. `floc=0`).

`fscale` similarly forces the scale to a particular value (e.g. `fscale=1`) .

It is necessary to use `floc` and/or `fscale` when computing MLEs if the distribution does not have a location and/or scale.

For example, the gamma distribution is defined using 2 parameters, often referred to as shape and scale.

In order to use ML to estimate parameters from a gamma, `floc=0` must be used.

► `dist.median`

Returns the median of the distribution. Generically, median is called using `dist.median(*args, loc=0, scale=1)`.

► `dist.mean`

Returns the mean of the distribution. Generically, mean is called using `dist.mean(*args, loc=0, scale=1)`.

► `dist.moment`

`nth` non-central moment evaluation of the distribution.

Generically, moment is called using `dist.moment(r, *args, loc=0, scale=1)` where `r` is the order of the moment to compute.

► `dist.var`

Returns the variance of the distribution. Generically, var is called using `dist.var(*args, loc=0, scale=1)`.

► `dist.std`

Returns the standard deviation of the distribution. Generically, std is called using `dist.std(*args, loc=0, scale=1)`.

Example

The gamma distribution is used as an example.

The gamma distribution takes 1 shape parameter a (a is the only element of `*args`), which is set to 2 in all examples.

```
>>> import scipy.stats as stats
>>> gamma = stats.gamma

>>> gamma.mean(2), gamma.median(2)
>>> gamma.std(2), gamma.var(2)
(2.0, 1.6783469900166608, 1.4142135623730951, 2.0)

>>> gamma.moment(2,2) gamma.
moment(1,2)**2 # Variance
```

```
>>> gamma.cdf(5, 2), gamma.pdf(5, 2)
(0.95957231800548726, 0.033689734995427337)
```

```
>>> gamma.ppf(.95957231800548726, 2)
5.00000000000000018
```

```
>>> log(gamma.pdf(5, 2)) gamma.
logpdf(5, 2)
0.0
```

```
>>> gamma.rvs(2, size=(2,2))
array([[ 1.83072394,  2.61422551],
       [ 1.31966169,  2.34600179]])

>>> gamma.fit(gamma.rvs(2, size=(1000)), floc = 0)
# a, 0, shape
(2.209958533078413, 0, 0.89187262845460313)
```