

# Contents

# 1 Data Structures

pandas introduces two new data structures to Python - Series and DataFrame, both of which are built on top of NumPy (this means it's fast).

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
pd.set_option('max_columns', 50)
```

## 1.1 Series

A Series is a one-dimensional object similar to an array, list, or column in a table. It will assign a labeled index to each item in the Series. By default, each item will receive an index label from 0 to N, where N is the length of the Series minus one.

```
# create a Series with an arbitrary list
s = pd.Series([7, 'Heisenberg', 3.14, -1789710578, 'Happy Eating!'])
s
```

```
0          7
1    Heisenberg
2         3.14
3   -1789710578
4    Happy Eating!
dtype: object
```

Alternatively, you can specify an index to use when creating the Series.

```
s = pd.Series([7, 'Heisenberg', 3.14, -1789710578, 'Happy Eating!'],
               index=['A', 'Z', 'C', 'Y', 'E'])
s
```

```
A          7
Z    Heisenberg
C          3.14
Y   -1789710578
E   Happy Eating!
dtype: object
```

The Series constructor can convert a dictionary as well, using the keys of the dictionary as its index.

```
In [4]: d = {'Chicago': 1000, 'New York': 1300, 'Portland': 900, 'San Francisco': 1100, 'Austin': 450, 'Boston': None}
cities = pd.Series(d)
cities
Out[4]: Austin 450 Boston NaN Chicago 1000 New York 1300 Portland 900 San Francisco 1100
dtype: float64
```

You can use the index to select specific items from the Series ...

```
In [5]: cities['Chicago']
Out[5]: 1000.0
In [6]: cities[['Chicago', 'Portland', 'San Francisco']]
Out[6]: Chicago 1000 Portland 900 San Francisco 1100
dtype: float64
```

Or you can use boolean indexing for selection.

```
In [7]: cities[cities < 1000]
Out[7]: Austin 450 Portland 900
dtype: float64
```

That last one might be a little weird, so let's make it more clear - `cities < 1000` returns a Series of True/False values, which we then pass to our Series `cities`, returning the corresponding True items.

```
In [8]: less_than_1000 = cities < 1000
print less_than_1000
print cities[less_than_1000]
bool
```

Austin 450 Portland 900 dtype: float64

You can also change the values in a Series on the fly.

In [9]: changing based on the index print 'Old value:', cities['Chicago']  
cities['Chicago'] = 1400 print 'New value:', cities['Chicago'] Old value: 1000.0  
New value: 1400.0

In [10]: changing values using boolean logic print cities[cities < 1000] print "  
cities[cities < 1000] = 750

print cities[cities < 1000] Austin 450 Portland 900 dtype: float64

Austin 750 Portland 750 dtype: float64

What if you aren't sure whether an item is in the Series? You can check using idiomatic Python.

In [11]: print 'Seattle' in cities print 'San Francisco' in cities False True

Mathematical operations can be done using scalars and functions.

In [12]: divide city values by 3 cities / 3 Out[12]: Austin 250.000000 Boston  
NaN Chicago 466.666667 New York 433.333333 Portland 250.000000 San Fran-  
cisco 366.666667 dtype: float64 In [13]: square city values np.square(cities)  
Out[13]: Austin 562500 Boston NaN Chicago 1960000 New York 1690000 Port-  
land 562500 San Francisco 1210000 dtype: float64 You can add two Series to-  
gether, which returns a union of the two Series with the addition occurring on  
the shared index values. Values on either Series that did not have a shared index  
will produce a NULL/NaN (not a number).

In [14]: print cities[['Chicago', 'New York', 'Portland']] print" print cities[['Austin',  
'New York']] print" print cities[['Chicago', 'New York', 'Portland']] + cities[['Austin',  
'New York']]

```
Chicago      1400
New York     1300
Portland      750
dtype: float64
```

```
Austin       750
New York     1300
dtype: float64
```

```
Austin       NaN
Chicago       NaN
New York     2600
Portland      NaN
dtype: float64
```

Notice that because Austin, Chicago, and Portland were not found in both Series, they were returned with NULL/NaN values.

NULL checking can be performed with `isnull` and `notnull`.

```
In [15]: returns a boolean series indicating which values aren't NULL
cities.notnull() Out[15]: Austin True Boston False Chicago True New York True
Portland True San Francisco True dtype: bool In [16]: use boolean logic to grab
the NULL cities print cities.isnull() print " print cities[cities.isnull()]
Austin False
Boston True Chicago False New York False Portland False San Francisco False
dtype: bool
```

```
Boston NaN dtype: float64
```

DataFrame

A DataFrame is a tabular data structure comprised of rows and columns, akin to a spreadsheet, database table, or R's `data.frame` object. You can also think of a DataFrame as a group of Series objects that share an index (the

column names).

For the rest of the tutorial, we'll be primarily working with DataFrames.

### Reading Data

To create a DataFrame out of common Python data structures, we can pass a dictionary of lists to the DataFrame constructor.

Using the columns parameter allows us to tell the constructor how we'd like the columns ordered. By default, the DataFrame constructor will order the columns alphabetically (though this isn't the case when reading from a file - more on that next).

```
In [17]: data = {'year': [2010, 2011, 2012, 2011, 2012, 2010, 2011, 2012],  
'team': ['Bears', 'Bears', 'Bears', 'Packers', 'Packers', 'Lions', 'Lions', 'Lions'],  
'wins': [11, 8, 10, 15, 11, 6, 10, 4], 'losses': [5, 8, 6, 1, 5, 10, 6, 12]}  
football = pd.DataFrame(data, columns=['year', 'team', 'wins', 'losses'])  
print football
```

```

year team wins losses 0 2010 Bears 11 5 1 2011 Bears 8 8 2 2012 Bears 10 6
3 2011 Packers 15 1 4 2012 Packers 11 5 5 2010 Lions 6 10 6 2011 Lions 10 6 7
2012 Lions 4 12

```

Much more often, you'll have a dataset you want to read into a DataFrame. Let's go through several common ways of doing so.

## CSV

Reading a CSV is as simple as calling the `read_csv` function. By default, the `read_csv` function

```
In [18]: /Users/greda/Dropbox/tutorials/pandas
```

```
In [19]: Source: baseball-reference.com/players/r/riverma01.shtml !head -n 5
mariano-rivera.csv Year, Age, Tm, Lg, W, L, W-L 1995, 25, NYY, AL, 5, 3, .625, 5.51, 19, 10, 2, 0, 0, 0
1996, 26, NYY, AL, 8, 3, .727, 2.09, 61, 0, 14, 0, 0, 5, 107.2, 73, 25, 25, 1, 34, 3, 130, 2, 0, 1, 425, 240, 0.994,
3MVP-12 1997, 27, NYY, AL, 6, 4, .600, 1.88, 66, 0, 56, 0, 0, 43, 71.2, 65, 17, 15, 5, 20, 6, 68, 0, 0, 2, 301, 2
25 1998, 28, NYY, AL, 3, 0, 1.000, 1.91, 54, 0, 49, 0, 0, 36, 61.1, 48, 13, 13, 3, 17, 1, 36, 1, 0, 0, 246, 233, 1.0
```

```
In [20]: from_csv = pd.read_csv('mariano - rivera.csv') from_csv.head()
```

Year	Age	Tm	Lg	W	L	W-L%	ERA	G	GS	GF	CG	SHO	SV	IP	H	R	ER	HR	BB	IBB	SO
HBP	BK	WP	BF	ERA+	WHIP	H/9	HR/9	BB/9	SO/9	SO/BB	Awards										
0	1995	25	NYN	AL	5	3	0.625	5.51	19	10	2	0	0	0	67.0	71					
43	41	11	30	0	51	2	1	0	301	84	1.507	9.5	1.5	4.0	6.9	1.70					
NaN																					
1	1996	26	NYN	AL	8	3	0.727	2.09	61	0	14	0	0	5	107.2	73					
25	25	1	34	3	130	2	0	1	425	240	0.994	6.1	0.1	2.8	10.9						
3.82 CYA-3MVP-12																					
2	1997	27	NYN	AL	6	4	0.600	1.88	66	0	56	0	0	43	71.2	65					
17	15	5	20	6	68	0	0	2	301	239	1.186	8.2	0.6	2.5	8.5	3.40					
ASMVP-25																					
3	1998	28	NYN	AL	3	0	1.000	1.91	54	0	49	0	0	36	61.1	48					
13	13	3	17	1	36	1	0	0	246	233	1.060	7.0	0.4	2.5	5.3	2.12					
NaN																					
4	1999	29	NYN	AL	4	3	0.571	1.83	66	0	63	0	0	45	69.0	43					
15	14	2	18	3	52	3	1	2	268	257	0.884	5.6	0.3	2.3	6.8	2.89					
ASCYA-3MVP-14																					

Our file had headers, which the function inferred upon reading in the file. Had we wanted to be more explicit, we could have passed `header=None` to the function along with a list of column names to use:

```
In [21]: Source: pro-football-reference.com/players/M/MannPe00/touchdowns/passing-
!head -n 5 peyton-passing-TDs-2012.csv 1,1,2012-09-09,DEN,,PIT,W 31-19,3,71,Demaryius
Thomas,Trail 7-13,Lead 14-13* 2,1,2012-09-09,DEN,,PIT,W 31-19,4,1,Jacob Tamme,Trail
14-19,Lead 22-19* 3,2,2012-09-17,DEN,@,ATL,L 21-27,2,17,Demaryius Thomas,Trail
0-20,Trail 7-20 4,3,2012-09-23,DEN,,HOU,L 25-31,4,38,Brandon Stokley,Trail 11-
31,Trail 18-31 5,3,2012-09-23,DEN,,HOU,L 25-31,4,6,Joel Dreessen,Trail 18-31,Trail
25-31
```

```
In [22]: cols = ['num', 'game', 'date', 'team', 'home_away', 'opponent', 'result', 'quarter',
pd.read_csv('peyton-passing-TDs-2012.csv', sep=',', header = None, names =
```



```
cols)no_headers.head()Out[22]: numgamedateteamhomeawayopponentresultquarterdist
09-09DENNaNPITW31-19371DemaryiusThomasTrail7-13Lead14-13*
1212012-09-09DENNaNPITW31-1941JacobTammeTrail14-19Lead22-
19*2322012-09-17DEN@ATLL21-27217DemaryiusThomasTrail0-
20Trail7-203432012-09-23DENNaNHOU25-31438BrandonStokleyTrail11-
31Trail18-314532012-09-23DENNaNHOU25-3146JoelDreessenTrail18-
31Trail25-31pandasvariousreaderfunctionshavemanyparametersallowingyoutodothi
```

There's also a set of writer functions for writing to a variety of formats (CSVs, HTML tables, JSON). They function exactly as you'd expect and are typically called `to_format`:

```
my_dataframe.to_csv('path_to_file.csv')
```

Take a look at the IO documentation to familiarize yourself with file reading/writing functionality.

## 1.2 Excel

Know who hates VBA? Me. I bet you do, too. Thankfully, pandas allows you to read and write Excel files, so you can easily read from Excel, write your code in Python, and then write back out to Excel - no need for VBA.

Reading Excel files requires the `xlrd` library. You can install it via `pip` (`pip install xlrd`).

Let's first write a `DataFrame` to Excel.

```
# this is the DataFrame we created from a dictionary earlier
print football.head()
```

year team wins losses 0 2010 Bears 11 5 1 2011 Bears 8 8 2 2012 Bears 10 6  
 3 2011 Packers 15 1 4 2012 Packers 11 5

```
# since our index on the football DataFrame is meaningless, let's not wr
football.to_excel('football.xlsx', index=False)

!ls -l *.xlsx
-rw-r--r--  1 greda  staff  5618 Oct 26 00:44 football.xlsx

# delete the DataFrame
del football

# read from Excel
football = pd.read_excel('football.xlsx', 'sheet1')
print football
```

	year	team	wins	losses
0	2010	Bears	11	5
1	2011	Bears	8	8
2	2012	Bears	10	6
3	2011	Packers	15	1
4	2012	Packers	11	5
5	2010	Lions	6	10
6	2011	Lions	10	6
7	2012	Lions	4	12

### 1.3 Databases

pandas also has some support for reading/writing DataFrames directly from/to a database [docs]. You'll typically just need to pass a connection object to the `read_frame` or `write_frame` functions within the `pandas.io` module.

Note that `write_frame` executes as a series of INSERT INTO statements and thus trades speed for simplicity. If you're writing a large DataFrame to a database, it might be quicker to write the DataFrame to CSV and load that directly using the database's file import arguments.

```
from pandas.io import sql
import sqlite3

conn = sqlite3.connect('/Users/greda/Dropbox/gregreda.com/_code/towed')
query = "SELECT * FROM towed WHERE make = 'FORD';"

results = sql.read_frame(query, con=conn)
print results.head()
```

	tow_date	make	style	model	color	plate	state	towed_address
0	01/19/2013	FORD	LL		RED	N786361	IL	400 E. Lower Wacker
1	01/19/2013	FORD	4D		GRN	L307211	IL	701 N. Sacramento
2	01/19/2013	FORD	4D		GRY	P576738	IL	701 N. Sacramento
3	01/19/2013	FORD	LL		BLK	N155890	IL	10300 S. Doty
4	01/19/2013	FORD	LL		TAN	H953638	IL	10300 S. Doty

	phone	inventory
0	(312) 744-7550	877040
1	(773) 265-7605	6738005

2	(773) 265-7605	6738001
3	(773) 568-8495	2699210
4	(773) 568-8495	2699209