

1 Useful Array and Matrix Functions

Many functions operate exclusively on array inputs, including functions which are mathematical in nature, for example computing the eigenvalues and eigenvectors and functions for manipulating the elements of an array.

1.1 Views

Views are computationally efficient methods to produce objects of one type which behave as other objects of another type without copying data. For example, an array `x` can always be converted to a matrix using `matrix(x)`, which will copy the elements in `x`. View “fakes” the call to `matrix` and only inserts a thin layer so that `x` viewed as a matrix behaves like a matrix.

1.1.1 view

`view` can be used to produce a representation of an array, matrix or recarray as another type without copying the data. Using `view` is faster than copying data into a new class.

```
>>> x = arange(5)
>>> type(x)
numpy.ndarray
>>> x.view(matrix)
matrix([[0, 1, 2, 3, 4]])
>>> x.view(recarray)
rec.array([0, 1, 2, 3, 4])
```

1.1.2 asmatrix, mat

`asmatrix` and `mat` can be used to view an array as a matrix. This view is useful since matrix views will use matrix multiplication by default.

```
>>> x = array([[1,2],[3,4]])
>>> x * x # Elementbyelement
array([[ 1,  4],
       [ 9, 16]])
>>> mat(x) * mat(x) # Matrix multiplication
matrix([[ 7, 10],
        [15, 22]])
```

Both commands are equivalent to using `view(matrix)`.

1.1.3 asarray

`asarray` work in a similar matter as `asmatrix`, only that the view produced is that of `ndarray`. Calling `asarray` is equivalent to using `view(ndarray)`

1.2 Shape Information and Transformation

1.2.1 shape

`shape` returns the size of all dimensions of an array or matrix as a tuple. `shape` can be called as a function or an attribute. `shape` can also be used to reshape an array by entering a tuple of sizes. Additionally, the new shape can contain -1 which indicates to expand along this dimension to satisfy the constraint that the number of elements cannot change.

```
>>> x = randn(4,3)
>>> x.shape
(4L, 3L)
>>> shape(x)
(4L, 3L)
>>> M,N = shape(x)
>>> x.shape = 3,4
>>> x.shape
(3L, 4L)
>>> x.shape = 6,-1
>>> x.shape
(6L, 2L)
```

1.2.2 reshape

`reshape` transforms an array with one set of dimensions and to one with a different set, preserving the number of elements. Arrays with dimensions M by N can be reshaped into an array with dimensions K by L as long as $M \cdot N = K \cdot L$. The most useful call to `reshape` switches an array into a vector or vice versa.

```
>>> x = array([[1,2],[3,4]])
>>> y = reshape(x,(4,1))
>>> y
```

```
array([[1],
       [2],
       [3],
       [4]])
>>> z=reshape(y,(1,4))
>>> z
array([[1, 2, 3, 4]])
>>> w = reshape(z,(2,2))
array([[1, 2],
       [3, 4]])
```

The crucial implementation detail of reshape is that arrays are stored using row-major notation. Elements in arrays are counted first across rows and then then down columns. reshape will place elements of the old array into the same position in the new array and so after calling reshape, $x(1) = y(1)$, $x(2) = y(2)$, and so on.

1.2.3 size

size returns the total number of elements in an array or matrix. **size** can be used as a function or an attribute.

```
>>> x = randn(4,3)
>>> size(x)
12
>>> x.size
12
```

1.2.4 ndim

ndim returns the size of all dimensions or an array or matrix as a tuple. **ndim** can be used as a function or an attribute .

```
>>> x = randn(4,3)
>>> ndim(x)
2
>>> x.ndim
2
```

1.2.5 tile

`tile`, along with `reshape`, are two of the most useful non-mathematical functions. `tile` replicates an array according to a specified size vector. To understand how `tile` functions, imagine forming an array composed of blocks. The generic form of `tile` is `tile(X, (M, N))` where `X` is the array to be replicated, `M` is the number of rows in the new block array, and `N` is the number of columns in the new block array.

```
>>> x = array([[1,2],[3,4]])
>>> z = hstack((x,x,x))
>>> y = vstack((z,z))
```

However, `tile` provides a much easier method to construct `y`

```
>>> w = tile(x,(2,3))
>>> y w
array([[0, 0, 0, 0, 0, 0],
       [0, 0, 0, 0, 0, 0],
       [0, 0, 0, 0, 0, 0],
       [0, 0, 0, 0, 0, 0]])
```

`tile` has two clear advantages over manual allocation:

- First, `tile` can be executed using parameters determined at run-time, such as the number of explanatory variables in a model,
- second `tile` can be used for arbitrary dimensions. Manual array construction becomes tedious and error prone with as few as 3 rows and columns.

`repeat` is a related function which copies data in a less useful manner.

1.2.6 ravel

`ravel` returns a flattened view (1-dimensional) of an array or matrix. `ravel` does not copy the underlying data (when possible), and so it is very fast.

```
>>> x = array([[1,2],[3,4]])
>>> x
```

```

array([[ 1, 2],
       [ 3, 4]])
>>> x.ravel()
array([1, 2, 3, 4])
>>> x.T.ravel()
array([1, 3, 2, 4])

```

1.2.7 flatten

`flatten` works much like `ravel`, only that it copies the array when producing the flattened version.

1.2.8 flat

`flat` produces a `numpy.flatiter` object (flat iterator) which is an iterator over a flattened view of an array. Because it is an iterator, it is especially fast and memory friendly. `flat` can be used as an iterator in a for loop or with slicing notation.

```

>>> x = array([[1,2],[3,4]])
>>> x.flat
<numpy.flatiter at 0x6f569d0>
>>> x.flat[2]
3
>>> x.flat[1:4] = 1
>>> x
array([[ 1, 1],
       [ 1, 1]])

```

1.2.9 split, vsplit, hsplit

`vsplit` and `hsplit` split arrays and matrices vertically and horizontally, respectively. Both can be used to split an array into `n` equal parts or into arbitrary segments, depending on the second argument. If scalar, the array is split into `n` equal sized parts. If a 1 dimensional array, the array is split using the elements of the array as break points. For example, if the array was `[2,5,8]`, the array would be split into 4 pieces using `[:2]`, `[2:5]`, `[5:8]` and `[8:]`. Both `vsplit` and `hsplit` are special cases of `split`, which can split along an arbitrary axis.

```

>>> x = reshape(arange(20),(4,5))
>>> y = vsplit(x,2)
>>> len(y)
2
>>> y[0]
array([[0, 1, 2, 3, 4],
       [5, 6, 7, 8, 9]])
>>> y = hsplit(x,[1,3])
>>> len(y)
3
>>> y[0]
array([[ 0],
       [ 5],
       [10],
       [15]])
>>> y[1]
array([[ 1, 2],
       [ 6, 7],
       [11, 12],
       [16, 17]])

```

1.2.10 delete

delete removes values from an array, and is similar to splitting an array, and then concatenating the values which are not deleted. The form of **delete** is **delete(x,rc,axis)** where rc are the row or column indices to delete, and axis is the axis to use (0 or 1 for a 2-dimensional array). If axis is omitted, **delete** operated on the flattened array.

```

>>> x = reshape(arange(20),(4,5))
>>> delete(x,1,0) # Same as x[[0,2,3]]
array([[ 0, 1, 2, 3, 4],
       [10, 11, 12, 13, 14],
       [15, 16, 17, 18, 19]])
>>> delete(x,[2,3],1) # Same as x[:,[0,1,4]]
array([[ 0, 1, 4],
       [ 5, 6, 9],
       [10, 11, 14],
       [15, 16, 19]])
>>> delete(x,[2,3]) # Same as hstack((x.flat[:2],x.flat[4:]))

```

```
array([ 0, 1, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18,
       19])
```

1.2.11 squeeze

squeeze removes *singleton* dimensions from an array, and can be called as a function or a method.

```
>>> x = ones((5,1,5,1))
>>> shape(x)
(5L, 1L, 5L, 1L)
>>> y = x.squeeze()
>>> shape(y)
(5L, 5L)
>>> y = squeeze(x)
```

1.2.12 fliplr, flipud

fliplr and **flipud** flip arrays in a left-to-right and up-to-down directions, respectively. **flipud** reverses the elements in a 1-dimensional array, and **flipud(x)** is identical to **x[::-1]**. **fliplr** cannot be used with 1-dimensional arrays.

```
>>> x = reshape(arange(4),(2,2))
>>> x
array([[0, 1],
       [2, 3]])
>>> fliplr(x)
array([[1, 0],
       [3, 2]])
>>> flipud(x)
array([[2, 3],
       [0, 1]])
```

1.2.13 diag

The behavior of **diag** differs depending depending on the form of the input.

- If the input is a square array, it will return a column vector containing the elements of the diagonal.
- If the input is an vector, it will return an array containing the elements of the vector along its diagonal.

Consider the following example:

```
>>> x = array([[1,2],[3,4]])
>>> x
array([[1, 2],
       [3, 4]])
>>> y = diag(x)
>>> y
array([1, 4])
>>> z = diag(y)
>>> z
array([[1, 0],
       [0, 4]])
```

1.2.14 triu, tril

`triu` and `tril` produce upper and lower triangular arrays, respectively.

```
>>> x = array([[1,2],[3,4]])
>>> triu(x)
array([[1, 2],
       [0, 4]])
>>> tril(x)
array([[1, 0],
       [3, 4]])
```


1.3 Some Useful Linear Algebra Functions

`det`

`det` computes the determinant of a square matrix or array.

```
>>> x = matrix([[1,.5],[.5,1]])
>>> det(x)
0.75
```

`solve`

`solve` solves the system $\mathbf{Ax}=\mathbf{b}$ when \mathbf{X} is square and invertible so that the solution is exact.

```
>>> A = array([[1.0,2.0,3.0],[3.0,3.0,4.0],[1.0,1.0,4.0]])
>>> b = array([[1.0],[2.0],[3.0]])
>>> solve(A,b)
array([[ 0.625],
       [1.125],
       [ 0.875]])
```

`eig`

`eig` computes the eigenvalues and eigenvectors of a square matrix. When used with one output, the eigenvalues and eigenvectors are returned as a tuple.

```
>>> x = matrix([[1,.5],[.5,1]])
>>> val,vec = eig(x)
>>> vec*diag(val)*vec.T
matrix([[ 1. , 0.5],
       [ 0.5, 1. ]])
```

`eigvals` can be used if only eigenvalues are needed.