1 Basic Functions and Numerical Indexing

1.1 Mathematics

`sum, cumsum`

`sum` sums elements in an array. By default, it will sum all elements in the array, and so the second argument is normally used to provide the axis to use – 0 to sum down columns, 1 to sum across rows. `cumsum` produces the cumulative sum of the values in the array, and is also usually used with the second argument to indicate the axis to use.

```
>>> x = randn(3,4)
>>> x
array([[0.08542071,
2.05598312,
2.1114733 , 0.7986635 ],
[0.17576066,
0.83327885, 0.64064119,
0.25631728],
[0.38226593,
1.09519101,
0.29416551, 0.03059909]])
>>> sum(x) # all elements
0.62339964288008698
>>> sum(x, 0) # Down rows, 4 elements
array([0.6434473
, 2.31789529,
1.76499762, 0.57294532])
>>> sum(x, 1) # Across columns, 3 elements
array([ 0.76873297, 0.23944028,
1.15269233])
>>> cumsum(x,0) # Down rows
array([[0.08542071,
2.05598312,
2.1114733 , 0.7986635 ],
[0.26118137,
1.22270427,
1.47083211, 0.54234622],
[0.6434473
, 2.31789529,
1.76499762, 0.57294532]])
```

sum and `cumsum` can both be used as function or as methods. When used as methods, the first input is the axis so that `sum(x,0)` is the same as `x.sum(0)`.

prod, cumprod

`prod` and `cumprod` behave similarly to `sum` and `cumsum` except that the product and cumulative product are returned. `prod` and `cumprod` can be called as function or methods.

diff

`diff` computes the finite difference of a vector (also array) and returnsn-1 an element vectorwhenused on an n element vector. diff operates on the last axis by default, and so diff(x) operates across columns and returns `x[:,1:size(x,1)]x[:,: size(x,1)1]` for a 2-dimensional array.

`diff` takes an optional keyword argument axis so that diff(x, axis=0) will operate across rows. `diff` can also be used to produce higher order differences (e.g. double difference).

```
>>> x= randn(3,4)
>>> x
array([[0.08542071,
2.05598312,
2.1114733 , 0.7986635 ],
[0.17576066,
0.83327885, 0.64064119,
0.25631728],
[0.38226593,
1.09519101,
0.29416551, 0.03059909]])
>>> diff(x) # Same as diff(x,1)
0.62339964288008698
>>> diff(x, axis=0)
array([[0.09033996,
2.88926197, 2.75211449,
1.05498078],
[0.20650526,
1.92846986,
0.9348067 , 0.28691637]])
>>> diff(x, 2, axis=0) # Double difference, columnbycolumn
array([[0.11616531,
4.81773183,
3.68692119, 1.34189715]])
```

exp

`exp` returns the element-by-element exponential (e^x) for an array.

log

`log` returns the element-by-element natural logarithm ($ln(x)$) for an array.

log10

`log10` returns the element-by-element base-10 logarithm (`log10 (x )`) for an array.

1.2 Rounding

around, round

around rounds to the nearest integer, or to a particular decimal place when called with two arguments.

```
%----------------%
\begin{framed}
\begin{verbatim}
>>> x = randn(3)
array([ 0.60675173, 0.3361189
, 0.56688485])
>>> around(x)
array([ 1., 0., 1.])
>>> around(x, 2)
array([ 0.61, 0.34,
0.57])
```

around can also be used as a method on an ndarray – except that the method is named round. For example, `x.round(2)` is identical to `around(x, 2)`. The change of names is needed to avoid conflicting with the Python built-in function round.

floor

`floor` rounds to the next smallest integer.

```
>>> x = randn(3)
array([ 0.60675173, 0.3361189
```

```
, 0.56688485])
>>> floor(x)
array([ 0., 1.,
1.])
```

ceil

`ceil` rounds to the next largest integer.

```
>>> x = randn(3)
array([ 0.60675173, 0.3361189
, 0.56688485])
>>> ceil(x)
array([ 1., 0.,
0.])
```

Note that the values returned are still floating points and so 0. is the same as 0..

1.3   Generating Arrays and Matrices

linspace

`linspace(l,u,n)` generates a set of n points uniformly spaced between l, a lower bound (inclusive) and u, an upper bound (inclusive).

```
>>> x = linspace(0, 10, 11)
>>> x
array([ 0., 1., 2., 3., 4., 5., 6., 7., 8., 9., 10.])
```

logspace

`logspace(l,u,n)` produces a set of logarithmically spaced points between 10l and 10u . It is identical to 10**linspace(l,u,n).

arange

`arange(l,u,s)` produces a set of points spaced by s between l, a lower bound (inclusive) and u, an upper bound (exclusive). arange can be used with a single parameter, so that

arange(n) is equivalent to arange(0,n,1). Note that arange will return integer data type if all inputs are integer.

```
>>> x = arange(11)
array([ 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10])
>>> x = arange(11.0)
array([ 0., 1., 2., 3., 4., 5., 6., 7., 8., 9., 10.])
>>> x = arange(4, 10, 1.25)
array([ 4. , 5.25, 6.5 , 7.75, 9. ])
```

1.3.1 `meshgrid`

`meshgrid` broadcasts two vectors to produce two 2-dimensional arrays, and is a useful function when plotting 3-dimensional functions.

```
>>> x = arange(5)
>>> y = arange(3)
>>> X,Y = meshgrid(x,y)
>>> X
array([[0, 1, 2, 3, 4],
[0, 1, 2, 3, 4],
[0, 1, 2, 3, 4]])
>>> Y
array([[0, 0, 0, 0, 0],
[1, 1, 1, 1, 1],
[2, 2, 2, 2, 2]])
```

1.3.2 `r_`

`r_` is a convenience function which generates 1-dimensional arrays from slice notation. While `r_` is highly flexible, the most common use it r_[start : end : stepOrCount] where start and end are the start and end points, and stepOrCount can be either a step size, if a real value, or a count, if complex.

```
>>> r_[0:10:1] # arange equiv
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
>>> r_[0:10:.5] # arange equiv
array([ 0. , 0.5, 1. , 1.5, 2. , 2.5, 3. , 3.5, 4. , 4.5, 5. ,
5.5, 6. , 6.5, 7. , 7.5, 8. , 8.5, 9. , 9.5])
>>> r_[0:10:5j] # linspace equiv, includes end point
array([ 0. , 2.5, 5. , 7.5, 10. ])
```

r_ can also be used to concatenate slices using commas to separate slice notation blocks.

```
>>> r_[0:2, 7:11, 1:4]
array([ 0, 1, 7, 8, 9, 10, 1, 2, 3])
```

Note that r_ is not a function and that is used with [].

1.3.3   c_

c_ is virtually identical to r_ except that column arrays are generates, which are 2-dimensional (second dimension has size 1).

```
>>> c_[0:5:2]
array([[0],
[2],
[4]])
>>> c_[1:5:4j]
array([[ 1. ],
[ 2.33333333],
[ 3.66666667],
[ 5. ]])
```

c_, like r_, is not a function and is used with [].

1.3.4   ix_

ix_(a,b) constructs an n-dimensional open mesh from n 1-dimensional lists or arrays. The output of ix_ is an n-element tuple containing 1-dimensional arrays. The primary use of ix_ is to simplify selecting slabs inside a matrix. Slicing can also be used to select elements froman array as long as the slice pattern is regular.

`ix_` is particularly useful for selecting elements froman array using indices which are not regularly spaced, as in the final example.

```
>>> x = reshape(arange(25.0),(5,5))
>>> x
array([[ 0., 1., 2., 3., 4.],
[ 5., 6., 7., 8., 9.],
[ 10., 11., 12., 13., 14.],
[ 15., 16., 17., 18., 19.],
[ 20., 21., 22., 23., 24.]])
>>> x[ix_([2,3],[0,1,2])] # Rows 2 & 3, cols 0, 1 and 2
array([[ 10., 11., 12.],
[ 15., 16., 17.]])
>>> x[2:4,:3] # Same, standard slice
array([[ 10., 11., 12.],
[ 15., 16., 17.]])
>>> x[ix_([0,3],[0,1,4])] # No slice equiv
```

1.3.5 mgrid

`mgrid` is very similar to meshgrid but behaves like `r_` and `c_` in that it takes slices as input, and uses a real valued variable to denote step size and complex to denote number of values. The output is an $n+1$ dimensional vector where the first index of the output indexes the meshes.

```
>>> mgrid[0:3,0:2:.5]
>>>
>>> mgrid[0:3:3j,0:2:5j]
>>>
```

1.3.6 ogrid

`ogrid` is identical to mgrid except that the arrays returned are always 1-dimensional. `ogrid` output is generally more appropriate for looping code, while `mgrid` is usually more appropriate for vectorized code. When the size of the arrays is large, then `ogrid` uses much less memory.

```
>>> ogrid[0:3,0:2:.5]
[array([[ 0.],
[ 1.],
[ 2.]]), array([[ 0. , 0.5, 1. , 1.5]])]
>>> ogrid[0:3:3j,0:2:5j]
[array([[ 0. ],
[ 1.5],
[ 3. ]]),
array([[ 0. , 0.5, 1. , 1.5, 2. ]])]
```